

EXECUTIVE REPORT · THE 35B SUITE · RECORDED 8 OCTOBER 2026 · COMPANION TO THE RE-RECORD REPORT

The same bank, a model a fifth of the size

Qwen3.6-35B on the live tier, and how it compares with the 122B

Model under test: Qwen3.6-35B-A3B-NVFP4 on the same two DGX Sparks; same books, prompts and rules
Ten live designs, 1,890 cells, 16,133 recorded calls, every design performed twice, in 95 minutes (the 122B: 317)
Compared design by design with the 122B re-record of 7–8 October; every figure replays offline in CI

PREPARED BY

Andrew Moore, Axiom Verity

Prepared for executive review. Internal. Evidence about a synthetic bank under one sample of each model;
not evidence about any real book or any frontier model.

Contents

Contents

Executive summary

1. What was run

1.1 The smoke

1.2 The run

2. The 35B's results

2.1 Agreement with the rule

2.2 Controls and the attack scenarios

2.3 Where the 35B fails: it does not call the tool

2.4 Reliability

2.5 The customer who answers back

3. The 35B beside the 122B

3.1 Where the 35B matches the 122B

3.2 Where the 35B does better

3.3 Where the 35B does worse

3.4 What the comparison adds to the record

4. Limits

5. Conclusions

6. Recommended next steps

Appendix A. Reproduction and artefacts

Appendix B. Glossary

2

3

4

4

5

6

6

6

7

8

8

9

10

10

10

12

12

12

13

14

14

Executive summary

On 8 October 2026 the ten live designs of the Craft A Bot synthetic bank ("Robo Bank") were re-performed with **Qwen3.6-35B-A3B** in the bot's seat, in place of the Qwen3.5-122B used for the 7–8 October re-record. Everything else was held constant: the same books, prompts, rules, reply limit (2,048 tokens), temperature (0) and timeout (180 s), on the same two DGX Sparks. Every design was performed twice. The run took 95 minutes against the 122B's 317, and every recording verifies against its own store and replays exactly in CI. This report presents the 35B's performance on its own terms, then compares it design by design with the 122B. It is a companion to the re-record report of 8 October, not a replacement for it.

94–100%

agreement with the rule on every desk once the rule is on the case file; three desks identical to the 122B

94 / 81

disputes decision matches the rule, 35B against 122B: the smaller model scores higher

40%

of advice bot cells lost: the 35B answers in prose where the desk needs a tool call

3.3×

faster: 95 minutes for 1,890 cells against 317 for 1,792, with every design performed twice

What the committee should take from this

- **Rule visibility generalises.** The 122B's headline, that seeing the rule was the whole difference, is not a property of the large model. With each desk's rule on its case file, the 35B reads 100% on onboarding, lending and collections, exactly as the 122B does, in a fifth of the time. A model a fifth of the size uses a rule it can see.
- **A larger model is not uniformly more accurate.** On disputes the 35B agrees with the rule in 94% of items (86–99%) against the 122B's 81% (69–93%). The model that misclassified about one dispute in five is not the smaller one. One design, 40 items, overlapping intervals: not a ranking, but a reason not to assume the large model is the safe baseline.
- **The 35B's failure mode is tool-call discipline, not judgement.** On the desks where the bot must speak through a tool to keep a conversation going, it often replies in plain prose instead: advice loses 40% of its bot cells, fraud 22%, collections 14%. In the preceding smoke, 77% of advice replies were text with no call. The prompts were tuned against the 122B, so this is a finding about the prompts' fit to a smaller model as much as about the model.
- **One control moved the wrong way.** On fraud, the policy-card stack takes alert decision from 94% to 81%, the only comparison in either suite where a control lowers a measure; the register's verdict is *not-supported*. The stack arm carries the stalls (a refused report re-filed, an alert re-opened), and that is as far as the data allows.
- **The 35B is less repeatable where it stalls.** Consistency across two performances is 100% on five desks but 93% on disputes and complaints and 89% on fraud (67% in one arm). pass² drops to 89–93% on those desks against the 122B's 96–98%.
- **The attack scenarios are at the same ceiling.** Unaided, the 35B resisted every attack exactly as the 122B did, so no agent-security component has a safety effect to measure on either model.

Read every figure as a measurement of this synthetic bank from one sample of each model at temperature 0, performed twice. The prompts were tuned against the 122B. It is not a statement about either model in general, nor about a real bank, a real customer or a regulator's rules.

Recommended next steps, in priority order

- Decide whether the engine tolerates a prose reply where a tool call is expected, or whether each model gets its own wording of the stage goals; the 35B's lost cells are almost entirely this.
- Read the fraud stories for the 35B's stack arm before any conclusion about the SAR card; the *not-supported* verdict rests on 27 items.
- Use the 35B as the "weaker model" the re-record report asked for, on the desks where it errs (disputes, fraud, complaints), so that controls have something to catch; the 122B sits at a ceiling there.
- Carry the comparison table into the Control Effectiveness Register as a second live column, with pass^k per model.
- A 35B-tuned run, as a separate experiment, to separate capability from prompt fit.

1. What was run

The suite is the one described in the re-record report: ten live designs derived from the reference experiments, with the brain factor replaced by a live model. The only change is the cartridge: *dgx-spark/quick-qwen*, serving Qwen3.6-35B-A3B-NVFP4, a mixture-of-experts model with about 3 billion active parameters, on both Sparks in the *fast-pair* pattern (chat mode, 262k context) rather than the 122B's *reasoning-pair*. Switching the Sparks between the two modes takes about seven minutes each way and was the plan's reserved owner decision.

	122B SUITE (7–8 OCTOBER)	35B SUITE (8 OCTOBER)
Model / cartridge	Qwen3.5-122B-A10B-NVFP4 / giant-qwen	Qwen3.6-35B-A3B-NVFP4 / quick-qwen
Spark pattern	reasoning-pair (puzzle mode)	fast-pair (chat mode, 262k context)
Designs, books, evaluators, factors, prompts	the ten live designs	the same ten, from the same base designs
Reply limit, temperature, timeout, concurrency	2,048 tokens; 0; 180 s; 16 cells	the same
Performed	twice for eight designs, once for the two servicing designs	twice for all ten
Cells / calls / wall time	1,792 / 14,469 / 317 min	1,890 / 16,133 / 95 min
Evidence	docs/evidence/live/	docs/evidence/live-35b/

Table 1. What is the same and what differs between the two suites (RUNBOOK.md).

1.1 The smoke

Before the run, 46 items across nine designs were recorded once against the 35B. The smoke predicted the stalls of the full run to within a few points, and it is where the failure mode was first seen: a large share of replies on some desks were prose with no tool call.




Figure 1. The smoke: share of replies that were text with no tool call, by design. The three desks that are a chain of lookups and decisions were at 0-3%; the desks that need a say action to keep a conversation going were not.

1.2 The run

DESIGN	PERFORMED	CELLS	CALLS	WALL TIME
onboarding	2×	132	780	3 min
disputes	2×	160	948	5 min
complaints	2×	220	582	3 min
servicing	2×	132	690	3 min
servicing, live customer	2×	64	690	3 min
controls-live (nine attack scenarios)	2×	162	1,681	11 min
collections	2×	148	1,889	9 min
lending	2×	408	2,564	15 min
advice	2×	248	3,879	32 min
fraud	2×	216	2,430	12 min
Total		1,890	16,133	95 min

Table 2. The 35B suite (FINDINGS.md §1). Every recording verified against its live store and reproduced by live-check --suite quick.

2. The 35B's results

2.1 Agreement with the rule

DESK	WHAT IS MEASURED	LIVE, NO STACK (95% INTERVAL, N)	ASSUMED 90% INSIDE THE INTERVAL?
Lending	decision matches the rule	100% (93–100%, n 51)	no
Servicing	the caller's need is met	100% (90–100%, n 33)	yes
Disputes	decision matches the rule	94% (86–99%, n 40)	yes
Collections	repayment plan matches the rule	100% (90–100%, n 36)	no
Onboarding	decision matches the rule	100% (90–100%, n 33)	yes
Complaints	root cause named	96% (93–99%, n 55)	no
Fraud	alert decision is the right one	94% (87–100%, n 27)	yes
Advice	recommendation suits the customer	100% (89–100%, n 31)	yes

Table 3. The 35B against the fallible tier's assumed 90% (README.md). Three desks sit above the assumption and outside their intervals.

With each desk's rule on its case file, the 35B agrees with it in 94–100% of items on every desk. The decision metrics are scored over the items whose cells completed; the share of cells lost before completion is in section 2.3.

2.2 Controls and the attack scenarios

- **Disputes.** Decision matches the rule 94% without the stack and 98% with it; reimbursed within the limit 99% and 100%. Both inside their intervals: *inconclusive*. The 35B was not tempted by the planted notes, so the *Within the limit* card, which took the 122B from 94% to 100%, has almost nothing to catch here.
- **Fraud.** Alert decision 94% without the stack and **81% with it**: the register's verdict is *not-supported*, and it is the only comparison in either suite where a control lowers a measure. Errors are spread across all four arms (10, 11, 11 and 15 of 54 cells), so it is not simply the stack's extra losses, but the stack arm also carries the stalls (a refused report re-filed, an alert re-opened). That is the reading the data allows, no more.
- **Attack scenarios.** Unaided, the 35B resisted every attack exactly as the 122B did: it kept the secret, the key, the ball and the code; sent no false alert and no malformed give; 100% in every scenario. No component has a safety effect to measure; the verdict is *inconclusive*. The components that act do so on termination: *no-progress* takes ran-out-of-steps from 100% to 83%, *marking* to 94%, *peer-auth* to 89%.


Fraud, with and without the stack (n 27 items, two performances)

Figure 2. Fraud, alert decision, with and without the policy-card stack: the 122B (left) and the 35B (right).

2.3 Where the 35B fails: it does not call the tool

Complaints (4% of bot cells lost), collections (14%), fraud (22%) and above all advice (40%) lose cells because the 35B often replies in plain prose when the desk needs a tool call. The failure repeats within a cell: in one advice cell the bot sent the same greeting and question twenty times as text until the stage ceiling ended the cell. The desks that need a *say* action to keep a conversation going are where it fails; the desks that are a chain of lookups and decisions (lending, disputes, onboarding) are where it does not. The stalls cluster on exactly the stages the 122B was fixed to avoid by rewording, so this is a finding about the model's tool-call discipline on these prompts, left as they are by the owner's choice of an unchanged suite.




Figure 3. Bot cells that did not end in success, 122B against 35B. A lost cell has no evaluation for its later stages.

2.4 Reliability




Figure 4. The 35B over items: pass@1, pass^2 and consistency, every design performed twice.

The 35B's non-determinism is its own and is read from its trials, not assumed from the 122B's. On five desks the two performances agree on every item. Where the model stalls, they do not: consistency is 93% on disputes and complaints and 89% on fraud (67% in one arm), and pass^2, the figure a control is credited on, falls to 90%, 93% and 89% there.

2.5 The customer who answers back

Servicing with a live 35B customer across the desk: needs met 94% (84–100%, n 16) against 100% (90–100%, n 33) with the scripted visitor; 9% of cells lost against 1%. The intervals overlap, so this is no detectable difference at this n. The 122B seat lost no cells.

3. The 35B beside the 122B

Both suites run the same ten designs on the same books with the same prompts, at temperature 0, with one model in the bot's seat. The 122B performed eight designs twice and the two servicing designs once; the 35B performed every design twice, so the reliability columns compare where both have them. The table is generated from the two suites' committed results by *live-compare.mjs*.

DESIGN	MEASURED	122B	35B	PASS^2 122B / 35B	CELLS LOST 122B / 35B	WALL TIME 122B / 35B	TOKENS A CASE 122B / 35B
Onboarding	decision matches the rule	100% (90–100%)	100% (90–100%)	100% / 100%	2% / 0%	19 / 3 min	13,540 / 8,126
Disputes	decision matches the rule	81% (69–93%)	94% (86–99%)	80% / 90%	4% / 1%	16 / 5 min	8,082 / 9,065
Complaints	root cause named	99% (97–100%)	96% (93–99%)	98% / 93%	0% / 4%	10 / 3 min	3,880 / 3,872
Servicing	needs met	100% (90–100%)	100% (90–100%)	– / 100%	0% / 1%	6 / 3 min	7,824 / 9,543
Servicing, live customer	needs met	100% (81–100%)	94% (84–100%)	– / –	0% / 9%	5 / 3 min	7,707 / 9,138
Attack scenarios	each attack resisted	100%	100%	– / –	– / –	37 / 11 min	27,025 / 26,783
Collections	plan matches the rule	100% (91–100%)	100% (90–100%)	100% / 100%	2% / 14%	26 / 9 min	16,356 / 27,324
Lending	agreement with the rule	100% (93–100%)	100% (93–100%)	100% / 100%	3% / 3%	66 / 15 min	12,391 / 11,866
Advice	recommendation suits	100% (89–100%)	100% (89–100%)	100% / 100%	0% / 40%	67 / 32 min	42,147 / 63,625
Fraud	alert decision	96% (82–99%)	94% (87–100%)	96% / 89%	12% / 22%	65 / 12 min	33,390 / 26,339

Table 4. The two suites design by design (COMPARISON.md). The decision metrics are over the items whose cells completed.




Figure 5. Agreement with each desk's rule in the reference configuration, 122B (navy) and 35B (gold), with 95% intervals. The dashed line is the fallible tier's assumed 90%.

3.1 Where the 35B matches the 122B

On onboarding, lending and collections, where the desk's rule is on the case file and the journey is a short chain of tool calls, the 35B reads the same 100% as the 122B, pass² included, in a fifth of the time. Servicing and advice suitability are also 100% on both. The attack scenarios read identically. The 122B's central finding, that putting the rule where the model can see it was the whole difference, holds for a model a fifth of the size.

3.2 Where the 35B does better

On disputes the 35B agrees with the rule in 94% of items against the 122B's 81%, and pass² is 90% against 80%. The model that misclassified about one dispute in five is not the smaller one. The 35B was also not tempted by the planted notes: reimbursed-within-limit reads 99% without the stack, where the 122B read 94%. The intervals overlap (69–93% against 86–99%) and it is one design of 40 items, so this is not a ranking; it is a reason not to assume the large model is the safe baseline.

3.3 Where the 35B does worse

- **Cells lost.** Advice 0% → 40%, fraud 12% → 22%, collections 2% → 14%, the servicing seat 0% → 9%, complaints 0% → 4%. The 122B's losses were the model repeating a call or retrying a refused action; the 35B's are prose where a tool call was needed.
- **Reliability.** pass² on complaints 98% → 93%, fraud 96% → 89%; consistency on fraud 100% → 89%. Where the 35B completes, it is as repeatable as the 122B; where it stalls, it is not.
- **Fraud's stack.** The 122B's alert decision is 96% in both arms; the 35B's falls from 94% to 81% with the stack. The 122B's stack arm also carried the stalls (11% of cells stopped) without moving the measure; the 35B's does move it.

- **Cost where it stalls.** Tokens per case are similar on most designs and higher where the model stalls: advice 42,147 → 63,625, collections 16,356 → 27,324. Wall time is 95 minutes against 317 overall.




Figure 6. Wall time per design. The 35B suite performed every design twice; the 122B performed eight of ten twice.




Figure 7. Tokens per case, no stack, both models.

3.4 What the comparison adds to the record

- **Rule visibility generalises.** It is a property of the desk design, not of the large model.
- **The prompts are tuned to one model.** The 35B's stalls are not random: they cluster on the desks and stages where the bot must speak through a tool, exactly the dead ends the 122B was fixed to avoid by rewording. A smaller model needs either its own wording or an engine tolerance for a prose reply, which is a design decision the owner declined for this suite so that the comparison stayed like for like.
- **A larger model is not uniformly more accurate.** Disputes is the counter-example; it should temper any instinct to treat the larger model as the conservative choice without measuring.
- **The 35B is the weaker model the re-record report asked for, on some desks.** The 122B sits at a ceiling almost everywhere, so its controls have nothing to catch. The 35B errs on disputes, complaints and fraud, which is where a control's effect could be measured; on fraud the first such measurement went the wrong way.
- **Cost.** A fifth of the parameters, a third of the wall time, every design performed twice, and the same answers on the desks that are a chain of lookups and decisions.

For a risk committee. Two of the findings cut against common assumptions. First, that the larger model is the safe default: on one desk it was the less accurate one. Second, that a control which is inert on one model is inert on another: the same SAR card that left the 122B's fraud measure untouched lowered the 35B's. Model choice is part of the control's operating environment, and a control's evidence does not transfer between models without being re-measured.

4. Limits

- **One sample of each model** at temperature 0, performed twice. The 35B's non-determinism is read from its trials (consistency 89–100%), not assumed from the 122B's.
- **Prompts tuned for the 122B.** The suite measures the 35B's fit to those prompts as much as its capability. A 35B-tuned run would be a separate experiment.
- **Lost cells are unscored later.** A cell lost to a stall has no evaluation for its later stages, so a design with many lost cells (advice, collections, fraud) is scored on fewer completed items than its n ; the lost share is in the table.
- **Small n .** Fraud's book is 27 items, the servicing seat 16. Most designs are at a ceiling and the register reads them *untestable* or *inconclusive*; none is *supported*. The fraud *not-supported* verdict rests on 27 items.
- **Synthetic bank and rules.** The rules, the incidence and temptation shares are this repository's assumptions awaiting review.
- **The register does not read either live column yet.**

5. Conclusions

- **A model a fifth of the size uses a rule it can see.** On the desks that are a chain of lookups and decisions, the 35B matches the 122B at 100%, in a third of the wall time.
- **Size did not decide accuracy.** On disputes the smaller model was the more accurate; on everything else the two are within each other's intervals.

- **The 35B's weakness is tool-call discipline on prompts written for another model.** It answers in prose where the desk needs a tool call, and loses up to 40% of a design's cells to it. Fix the prompts or the engine, not the model, first.
- **Controls do not transfer between models without re-measurement.** The fraud SAR card was inert on the 122B and lowered the 35B's alert decision from 94% to 81%.
- **The 35B gives the programme what it lacked:** a model that errs on some desks, so that controls there have something to catch.



6. Recommended next steps

#	ACTION	WHY	HORIZON
1	Decide the engine's response to a prose reply where a tool call is expected: tolerate it (and score it), or give each model its own wording of the stage goals.	Almost all of the 35B's lost cells are this one behaviour	Immediate (design decision)
2	Read the fraud stories for the 35B's stack arm before any conclusion is drawn about the SAR card.	The not-supported verdict rests on 27 items and the stack arm carries the stalls	Immediate
3	Use the 35B as the weaker model on disputes, complaints and fraud , where it errs, so that the controls there have an effect to measure.	The 122B sits at a ceiling on those desks	Short term
4	Add the 35B as a second live column in the Control Effectiveness Register , with pass^k per model.	A control's evidence is model-specific; the register should say which model it was measured on	Short term
5	Run a 35B-tuned suite as a separate experiment , rewording the stage goals where it stalled.	Separates capability from prompt fit	Medium term
6	Harder books for the desks at a ceiling on both models (onboarding, lending, collections, advice).	Neither model is tempted by the current temptations	Medium term

Table 5. Recommended actions, in priority order.

Recommended decision. Keep both live columns as evidence, each labelled with its model; take the engine decision on prose replies before any further 35B run; and treat the fraud result as a question to investigate, not a finding to act on.

Appendix A. Reproduction and artefacts

WHAT	WHERE
This report's sources	docs/evidence/live-35b/FINDINGS.md (the narrative); COMPARISON.md (generated, both suites); README.md (the generated column); RUNBOOK.md
Per design: cassette, result, markdown, design, cells.json, stories	docs/evidence/live-35b/<design>-live/
The live runs' own stores	recordings/35b/ (local, gitignored); the smoke under recordings/smoke-quick/
The 122B suite this is compared with	docs/evidence/live/ (TEST-REPORT.md, RUNS-AND-FINDINGS.md, README.md)
Replay every 35B design offline, held to the committed result	node scripts/live-check.mjs --suite quick
Regenerate the column and the comparison	node scripts/live-column.mjs --suite quick --check; node scripts/live-compare.mjs --check
The run (needs the Sparks in fast-pair)	node scripts/live-record.mjs --suite quick --all --resume
Switch the Sparks	craftabot spark plan / up --pattern fast-pair; spark down restores the puzzle mode

Appendix B. Glossary

TERM	MEANING
Suite	The ten live designs run with one model in the bot's seat; the 122B suite and the 35B (quick) suite share designs, books and prompts.
Cell / item / performance	A cell is one case through one journey under one configuration; an item is a case; a performance is one trial of it with fresh model draws.
Cells lost	Bot cells that did not end in success: an error or a stack stop. The decision metrics are scored over the items whose cells completed.
pass@1 / pass ² / consistency	One performance passes / both pass (the figure for a control) / the two performances agree. Over items, never trials.
Stack / policy cards	The guard control under test: policy cards placed between the bot and its actions.
Temptation	A planted provocation in a book that a guard is meant to catch.
Say action	The tool call by which the bot speaks to the customer; a reply in plain prose outside it does not advance the stage.
Verdict	untestable (both sides at the same bound), inconclusive (intervals overlap), supported / not-supported (effect excludes zero, in the expected or the opposite direction).
fast-pair / reasoning-pair	The Sparks' serving patterns: both units in chat mode for the 35B; both in puzzle mode for the 122B.