

Running a bank's desks on a live model: runs and findings

Executive report on work packages WP168 and WP169

Model under test: Qwen3.5-122B-A10B-NVFP4, served on two NVIDIA DGX Sparks

Nine reference designs, ten recordings, 1,272 cells, 8,641 cassette entries, 3 h 32 min of recording

Recorded and analysed 6 October 2026

Contents

- Contents 2**
- Executive summary 3**
- 1. Purpose and context 4**
 - How to read the numbers 4
- 2. What was run 4**
 - The model and the hardware 5
 - How a run works 5
- 3. Findings 6**
 - 3.1 The live error rate against the assumed 90% 6
 - 3.2 Two controls with large adverse side-effects 7
 - 3.3 Desk-by-desk results 8
 - 3.4 Cells that ended in ERROR 8
 - 3.5 Customers who answer back (WP169) 9
 - 3.6 The live tier's own variance 9
- 4. Cost and footprint 10**
- 5. Problems found on the way, and their fixes 10**
- 6. Limits of the evidence 12**
- 7. Conclusions 12**
- 8. Recommended next steps 12**
- Appendix A. Where things are 14**
- Appendix B. Glossary 14**

Executive summary

On 6 October 2026 the Craft A Bot synthetic bank ("Robo Bank") was run for the first time with a live language model making the decisions on every desk, in place of the scripted and "fallible" stand-ins used until now. The model was Qwen3.5-122B-A10B-NVFP4, a model of the kind a bank could run entirely inside its own boundary, served on two NVIDIA DGX Sparks. Eight reference designs with a customer book were re-run (lending twice), and a ninth put a second live model across the desk as the customer. Every run is committed as a replayable "cassette", and the build pipeline now re-derives every published number from those cassettes with no network access.

4 of 8

desks where the assumed 90% agreement is outside the measured 95% interval

25%

root cause named on the complaints desk (16–38%), against 90% assumed

20%

complaints redress within bounds with the policy-card stack on (100% without it)

0%

advice data minimisation under relational context (100% under the case file)

What the committee should take from this

- **The bank's uniform error assumption does not survive contact with a live model.** The fallible tier assumed one decision in ten wrong on every desk. Lending (53%), disputes (68%), collections (72%) and complaints (25%) are far worse than that; onboarding and advice (100%) are better. One number per desk cannot stand in for a live model; the shape of its errors differs by desk. The cassettes now hold the evidence to replace the assumption with measured, desk-shaped error models.
- **Two controls produced large adverse side-effects on properties they were not aimed at.** The policy-card guard stack, the programme's principal control, left the complaints root-cause metric where it was but collapsed "redress within bounds" from 100% to 20% ($p \approx 10^{-13}$). Giving the advice bot a richer relational context bought nothing measurable on suitability and made it fail data minimisation on all 31 cases ($p \approx 10^{-15}$). Neither effect has yet been diagnosed.
- **No control has yet been shown to work.** Where policy cards were the treatment, the effect on the primary metric was not detectable at these sample sizes (lending p 0.51, collections p 1.0, complaints p 0.45, disputes identical). Every verdict is deliberately *inconclusive*, *untestable* or *not-supported*; none is *supported*.
- **Temperature 0 is not repeatable on this serving stack.** Recording lending twice with nothing changed gave the same aggregate (53% both times) but only 63% of identical prompts received the same call and 91% of cases the same verdict. Any control whose effect is smaller than that noise cannot be credited from one recording.
- **The tooling is now sound enough to be relied on.** Six defects found on the way (a reply cap that measured itself, a credential guard tripping on a placeholder key, oversized cassettes, a model id not reaching the agent, the live customer not being recorded, a false positive in the synthetic-data sweep) were each fixed, and CI replays every design from its cassette and requires exactly the committed numbers.

Read every figure as a measurement of this synthetic bank, at the size stated, under one sample from one model. It is evidence about a model a bank could run inside its own boundary. It is not evidence about a frontier model or about any real book, and the rules the bot is scored against are this repository's rules, not a regulator's.

Recommended next steps, in priority order

- Diagnose the two large adverse effects (complaints redress; advice data minimisation) from the committed stories before any further control is added to either desk.
- Replace the uniform 10% error assumption with desk-shaped error models derived from the cassettes (WP170), and fold the live column into the Control Effectiveness Register and Control Inventory.
- Triage the ERROR cells (complaints 40%, fraud 34%, advice 18%) and separate model, prompt and desk causes.
- Raise statistical power before crediting any control: larger books, a second recording for every design, and a minimum detectable difference that is set in advance.
- Extend the live tier to the scenario designs and to a second model, so that findings can be told apart from artefacts of one model and one serving stack.

1. Purpose and context

Craft A Bot is a simulator of a UK retail bank operated by AI agents. Each **desk** (lending, servicing, disputes, collections, onboarding, complaints, fraud, advice) is a workflow in which an agent, the **bot**, works through a **book** of synthetic customers and cases and is scored against the desk's rules. A **reference design** is an experiment on one desk: the factors to vary (for example whether a guard is in place, or who executes the decision), the metrics to score, and the statistics to report. A **cell** is one case run under one configuration.

Until this work the bot's **brain** was one of two stand-ins. The **scripted** brain follows the rules perfectly; the **fallible** brain follows them with an assumed error rate, set in *fs-bank*'s *ERROR_RATES* at one decision in ten on every desk. Those two columns are the evidence base behind the programme's Control Effectiveness Register and Control Inventory. The **live tier** (work package WP168) replaces the brain with a real language model; WP169 adds a live model as the *customer* as well, so the conversation across the desk is two models talking.

The question the live tier answers is therefore not "how good is the model" but "how far are the assumptions the bank has been testing its controls against from what a live model actually does, and do the controls still behave as expected when they meet one?"

How to read the numbers

- Every percentage is a proportion of cells with a 95% **Wilson interval** and the number of cells (n). Effect intervals are Newcombe intervals; p-values are from the design's own test.
- Effects are guard *none* against *policy-cards* unless another factor is named.
- Most designs are **underpowered** at these sizes: fewer than 30 cells on a side, or a minimum detectable difference far above the five points the design aimed for. "No difference" means *not detectable here*, never *no effect*.
- Every figure comes from the committed results (*experiment-result.json*, *cells.json*, *timings.json*) and is re-derived by CI from the cassette alone.

2. What was run

Eight reference designs with a book were re-run with the brain a live model, and lending was recorded twice to measure the tier's own variance. A ninth recording re-ran servicing with the customer also a live model. Book sizes differ because desks' incidences differ widely and each design had to record in one session: the fraud book draws nearly five alerts per customer, so six customers already give 108 cells.

Design	Book size	Cells	Cassette entries	Wall time	Stories kept
lending-stack	800	306	1,238	30 min	18
lending-stack (second recording, -b)	800	306	1,099	27 min	18
servicing-stack	200	66	260	9 min	4
disputes-stack	400	80	265	6 min	4
collections-stack	300	74	839	20 min	10
onboarding-stack	400	66	289	6 min	2
complaints-stack	200	110	302	7 min	5
fraud-stack	6	108	2,490	54 min	16
advice-context	1,200	124	1,502	46 min	8
servicing-stack with a live customer (-seat)	100	32	357	8 min	2
Total		1,272	8,641	3 h 32 min	87

Table 1. The ten recordings of 6 October 2026. Cells across the ten recordings total 1,272; across the nine distinct designs (lending counted once) 966.

The model and the hardware

The model was **Qwen3.5-122B-A10B-NVFP4** (cartridge *dgx-spark/giant-qwen*), a 122-billion-parameter mixture-of-experts model quantised to NVFP4, served on the builder's two NVIDIA DGX Sparks in the *reasoning-pair* pattern: both units serving the model in *puzzle* mode, eight streams each, with multi-token-prediction speculative decoding on. Recording ran sixteen cells at a time with load spread across both units. Temperature was 0 and every reply was capped at 1,024 tokens. This is a model a bank could run inside its own boundary; it is not a frontier model.

How a run works

- **Derive.** *live-designs.mjs* derives each live design from its base design: same template, factors and metrics; the brain factor removed; one live brain naming its cassette; a book of its own size; a budget for the live cells; a 1,024-token reply limit on every build.
- **Record.** *live-record.mjs* checks that the Sparks serve the cartridge, records with *craftabot record* at automatic concurrency, then replays the design from the cassette alone with egress set to none.
- **Commit.** The replay's result, markdown, per-case *cells.json* and a sample of runs as readable stories are committed beside the slimmed cassette (about 1 MB each).
- **Verify.** CI replays every live design from its cassette with no network and requires exactly the committed numbers (*live-check.mjs*), and regenerates the README's live column (*live-column.mjs*).

The result is an audit trail: every number in this report can be re-derived, offline, from a committed artefact, and the build fails if it cannot. That is the property a model-risk function would ask of any evidence a control's effectiveness rests on.

3. Findings

3.1 The live error rate against the assumed 90%

The fallible tier assumed an agreement of 90% on every desk. The live baseline, reference configuration with no guard, is shown below. Four of the eight desks fall far below the assumption; three sit on it; two sit above it at the ceiling.




Figure 1. Live baseline per desk with 95% Wilson intervals (n in brackets). The dashed line is the 90% the fallible tier assumed. Amber desks have the assumption outside their interval.

Desk	What is measured	Live (95% interval, n)	Assumed 90% inside the interval?
Lending	decision matches the rule	53% (40–66%, n 51)	no
Disputes	decision matches the rules	68% (52–80%, n 40)	no
Collections	repayment plan matches the rule	72% (56–84%, n 36)	no
Complaints	root cause named	25% (16–38%, n 55)	no
Servicing	caller's need met	94% (80–98%, n 33)	yes
Fraud	alert decision right	93% (77–98%, n 27)	yes
Onboarding	decision matches the rules	100% (90–100%, n 33)	yes
Advice	recommendation suits the customer	100% (89–100%, n 31)	yes

Table 2. The first finding in numbers.

The uniform 10% assumption is wrong in both directions: too optimistic by a wide margin for four desks and too pessimistic for onboarding and advice. A single number per desk cannot stand in for a live model; the shape of its errors differs by desk. This is the finding that the shaped error models planned in WP170 were provisional for, and the cassettes now hold the evidence to replace the assumptions.

What this does not show is why lending, disputes, collections and complaints are low. That was not diagnosed; the cases and stories are in each design's folder. The complaints figure has a history: WP155 once found the desk's truth and rule disagreeing with the register. Whether the 25% is the model, the desk, or the task as posed has not been examined.

3.2 Two controls with large adverse side-effects

The largest effects in the live tier are not improvements. In both cases the metric the treatment was aimed at did not move, and a different property the bank relies on collapsed.

Two properties that collapsed under the treatment (metric otherwise unchanged)




Figure 2. Left: complaints, redress within bounds, guard none against the policy-card stack (n 55 a side). Right: advice, data minimisation under the case-file context against the relational context (n 31 a side).

Complaints: the policy-card stack and redress

Complaint acknowledged and ombudsman disclosed were 100% in both arms. Root cause named moved from 25% to 31% (p 0.45), no detectable effect. But **redress within bounds** was 100% without the stack and **20% (12–32%) with it** ($p \approx 10^{-13}$). 44 of the design's 110 cells ended in ERROR. The policy-card stack is the treatment and the metric collapses with it. It could be the stack blocking a path the bot then fails to complete, or an interaction between the cards and how a live bot proposes redress. The complaints stories are the place to start, and this has not been done.

Advice: richer context and data minimisation

The advice design is a context ladder: the same bot given either a case file or a richer relational context. Suitability of the recommendation was 100% under both. **Data minimisation** was 100% under the case file and **0% (0–11%) under the relational context** ($p \approx 10^{-15}$): given more about the customer, the live bot failed to minimise on every one of 31 cases, while tokens per case rose from about 58k to 82k. On this sample, richer context bought nothing measurable on suitability and cost the minimisation property. Whether that holds for other models or prompts is not known.

Why this matters for a risk committee. Both effects are the pattern a control framework should fear most: an intervention that leaves its target metric unchanged while silently degrading a conduct or data-protection outcome. Had only the target metric been monitored, both would have been invisible. The live tier caught them because every design scores several properties at once.

3.3 Desk-by-desk results

Desk	Primary result (reference, no guard)	Control effect	Other properties	Outcomes (success / error / stopped)	Verdict
Lending (two recordings)	Agreement 53% (40–66%, n 51); over-approval 0% (0–7%)	Policy cards 47% (34–60%), p 0.51: none detectable. <i>rules-only</i> executor reads 100%, most likely the rule agreeing with itself (not checked). Person at the decision reads 53%, same as the bot alone	≈15k tokens/case; 17k with the stack (12.7k in the second recording)	279 / 13 / 14 (first); 286 / 18 / 2 (second)	not supported
Servicing	Needs met 94% (80–98%, n 33)	Stack changes nothing at this n	Disclosure recorded 100%; ≈11.8k tokens/case	62 / 2 / 2	inconclusive
Disputes	Decision matches the rules 68% (52–80%, n 40)	Stack identical to no stack, same numbers and tokens: the cards did not act on this book	Reimbursed within the limit 100%; ≈7.5k tokens/case	80 / 0 / 0	not supported
Collections	Plan matches the rule 72% (56–84%, n 36)	76% (59–87%) with the guard, p 1.0: none detectable	Vulnerability actioned 100%; 30.4k tokens/case without the stack, 23.3k with	60 / 8 / 6	not supported
Onboarding	Decision matches the rules 100% (90–100%, n 33)	Nothing for a control to improve	Tipping-off gate (hit-contained) 100%	66 / 0 / 0	untestable (ceiling)
Complaints	Root cause named 25% (16–38%, n 55)	31% (20–44%) with the stack, p 0.45: none detectable. Redress within bounds 100% → 20%	Acknowledged 100%; ombudsman disclosed 100%	44 of 110 ERROR	adverse effect, undiagnosed
Fraud	Alert decision right 93% (77–98%, n 27)	A person at the SAR stage reads identically	No tipping-off 100%; ≈107k tokens/case, 70k with the stack; 2,490 cassette entries	42 / 37 / 29	inconclusive
Advice (context ladder)	Suitable recommendation 100% (89–100%, n 31) under both contexts	Data minimisation 100% → 0% under the relational context	≈58k tokens/case, 82k with relational context	102 / 22 / 0	adverse effect, undiagnosed

Table 3. Each design's live column. "Stopped" is STOPPED_BY_GUARDRAIL.

3.4 Cells that ended in ERROR

An ERROR cell is a stage that did not produce the output its schema asked for, which a live model can do for reasons about the model, the prompt or the desk. They are scored as the evaluator sees them and have not been diagnosed. Three designs carry many; a sampled story per outcome is committed for each.




Figure 3. Share of cells ending ERROR, by design.

3.5 Customers who answer back (WP169)

Until WP169 a book's customer was either absent or the desk's scripted visitor. A book campaign whose counterpart is *live* now seats a second live model across the desk at each agent stage; each line it says becomes a *seat.said* event on the bot's trace. Servicing was re-run with 100 customers and the 122B as both bot and customer.

Customer	Needs met (95% interval, n)	Tokens per case	Cells
The desk's own scripted visitor	94% (80–98%, n 33)	11,775	66
A live customer	100% (81–100%, n 16)	8,735	32

Table 4. Servicing with a scripted and a live customer, reference configuration, no guard.

All 32 cells completed. The intervals overlap and the books differ in size, so this is **no detectable difference**, not evidence that a talking customer helps. What it demonstrates is that a two-model conversation runs end to end on the live tier and replays exactly from one cassette. One weakness is visible in the stories: the persona is drawn by cohort rather than by the request, so a customer may open about savings when the case is a lost card. The bot copes, but a persona matched to the request would be a better test.

3.6 The live tier's own variance

Lending was recorded twice with nothing changed between the recordings. A recording is one sample, so the difference between the two is the tier's own noise, the floor under any effect a control is credited with.

63%

of the 488 prompts both recordings asked were answered with the same call (same tool, same arguments)

8%

were answered with the same words as well

91%

of the 306 cases both recordings decided reached the same verdict

0.0 pts

difference in the headline figure: 53% (40–66%, n 51) in both

1,361 prompts were asked by only one recording, because a different earlier answer leads to a different later prompt. Finish reasons were all *tool_call* in the first recording and all but two in the second; none was cut off by

the length limit. **Temperature 0 is not repeatable on this serving stack:** batching and speculative decoding move the numerics with whatever else is in flight. The aggregate is stable here; individual calls and cases are not. Any control whose effect is smaller than this noise cannot be credited from one recording, which is exactly what the underpowered flags say.

4. Cost and footprint

The whole tier recorded in three and a half hours of wall time on two desktop-class units, which makes a repeat recording, or a second model, a matter of an afternoon rather than a procurement. Token cost per case varies by more than an order of magnitude across desks, driven by how many stages and tool calls a case takes.




Figure 4. Approximate tokens per case in the reference configuration. Fraud is the heaviest desk; the policy-card stack reduced tokens on fraud (107k to 70k) and collections (30.4k to 23.3k) and raised them on lending.




Figure 5. Wall time per recording, sixteen cells at a time across both Sparks.

5. Problems found on the way, and their fixes

Six defects cost time during the day. Each was fixed and each is a lesson about what a live tier needs that a scripted one does not.

What happened	Cause	Fix
35 of the first lending recording's 791 replies (4.4%) were cut off; agreement read 10 of 23	Stage agents inherited the starter's 256-token reply limit. The run measured the cap, not the model	Live designs set a 1,024-token limit; no reply in any later recording was cut off. The 256-token recording was discarded
28 minutes of lending recording discarded at the end	The credential guard refused to write: a 3-character placeholder for a guard key in .env matched ordinary words in model replies	A smoke variable shorter than 8 characters is treated as a placeholder, the key check says so, and a recorder now stops at the first response carrying a held credential
Cassettes were about 14 MB each	Raw wire chunks were recorded with every answer	Slim recording drops them (about 1 MB per cassette); session-visible responses are slimmed identically so replay matches
A live brain asked its provider for the model "mock"	The brain's cartridge id never reached the agent's brain slot	The campaign now applies it, with a test
With a live customer, 26 of 32 cells ended ERROR on replay (cassette miss)	The recorder only captured brains mapped to a cassette; the customer's brain id differed, so none of its answers were recorded	The customer shares the brain's id, so one recording writes both into one cassette
Synthetic-data sweep flagged Luhn-valid 16-digit floats in result JSON	Statistics, not card numbers	The live sweep now covers cassettes and stories only

Table 5. Defects found during the recordings of 6 October 2026.

Two of these deserve the committee's attention as general lessons. A reply cap inherited silently from a starter produced a plausible-looking but meaningless result (10 of 23) until the cut-off replies were counted; **a live tier must record finish reasons and treat any truncation as invalidating**. And a credential guard configured with a placeholder key produced a false positive that destroyed half an hour of evidence; **guards need their own smoke tests**, and the fix (stop at the first held credential rather than at the end) is the right failure mode.

6. Limits of the evidence

- **One sample, one model.** Each design is a single recording. The second lending recording shows the aggregate reproducing, but only for that design.
- **Small n.** Most effects are underpowered. The verdicts are *inconclusive*, *untestable* or *not-supported*, deliberately; none is *supported*.
- **Synthetic bank.** Every customer, case and record is generated; the rules the bot is scored against are this repository's rules, not a regulator's.
- **ERROR cells are not diagnosed.** Complaints 44 of 110, fraud 37 of 108, advice 22 of 124. Each is a stage that did not produce the output its schema asked for; the cause may be the model, the prompt or the desk.
- **The register does not read these results yet.** The Control Effectiveness Register and the Control Inventory still read the scripted and fallible columns; the live results sit beside them, not folded in.
- **Not recorded.** The scenario designs (*controls*, *gate-presets*) were not run live, and no edition serves a cassette, so the Workshop's live column is empty.
- **Not separately checked.** The *rules-only* executor's 100% on lending is most likely the rule agreeing with itself; and why a person at the decision reads the same as the bot alone was not examined.

7. Conclusions

- **The programme's error assumption is falsified as a uniform number.** Four desks are far worse than the assumed 90% and two are better. Any control-effectiveness conclusion that rested on the fallible tier's 10% should be treated as provisional until it is re-derived against desk-shaped, measured error models.
- **Controls have not been shown effective, and two have been shown to have adverse side-effects.** The policy-card stack and the relational context each left their target metric unchanged while collapsing a conduct or data-protection property. These are the first findings in the programme with p-values that survive the tier's noise, and they point the wrong way.
- **Multi-property scoring is what caught them.** The design discipline of scoring several properties on every desk, not only the one a control is aimed at, is the reason the side-effects are visible at all. It should be non-negotiable in any real-bank equivalent.
- **Reproducibility must be measured, not assumed.** Temperature 0 on a batched, speculating serving stack is not deterministic. The second recording gives the tier a noise floor; every future claim of effect needs to clear it.
- **The infrastructure is ready.** Derive, record, commit, verify; CI replaying every number offline; two-model conversations recorded into one cassette. The remaining work is statistical power and diagnosis, not plumbing.

8. Recommended next steps

#	Action	Why	Horizon
1	Diagnose the complaints redress collapse from the committed complaints stories: is the stack blocking a path the bot then fails to complete, or does it change how a live bot proposes redress?	Largest effect in the tier; a control degrading a conduct outcome	Immediate
2	Diagnose the advice data-minimisation failure under relational context, and test whether it holds for a second model or prompt.	Data-protection property lost on 31 of 31 cases	Immediate

#	Action	Why	Horizon
3	Triage ERROR cells (complaints, fraud, advice) into model, prompt and desk causes; decide whether each is scored as a failure or excluded.	Up to 40% of cells in a design; they move every headline	Immediate
4	Investigate the four low desks (lending, disputes, collections, complaints), starting with whether the complaints truth, rule and register agree (WP155).	Separates model weakness from desk or task defects	Short term
5	Replace the uniform 10% assumption with desk-shaped error models fitted to the cassettes (WP170); re-run the fallible tier against them.	The first finding; the fallible tier currently tests controls against the wrong adversary	Short term
6	Fold the live column into the Control Effectiveness Register and Control Inventory , with the verdict vocabulary (supported / inconclusive / untestable / not-supported / adverse) carried through.	The register is what the governance lenses read	Short term
7	Set power targets in advance : book sizes for at least 30 cells a side and a minimum detectable difference of 5 points; record every design at least twice.	Nothing can be credited below the measured noise floor	Short term
8	Record the scenario designs (controls, gate-presets) live and serve a cassette in an edition so the Workshop's live column is populated.	Completes the tier	Medium term
9	Add a second model (another local model, and a frontier model via a hosted key) to tell model-specific findings from general ones.	One sample from one model is the tier's chief limit	Medium term
10	Match live-customer personas to the request , and seat live customers on the desks beyond servicing.	Makes the two-model conversation a fair test	Medium term
11	Verify the rules-only 100% and the bot-with-a-person result on lending, so executor comparisons mean what they appear to.	Avoids a misleading comparator in the register	Medium term

Table 6. Recommended actions, in priority order.

Recommended decision. Treat the live tier as the programme's primary evidence base from this point, keep the fallible tier only as a calibrated stand-in derived from it, and do not add further controls to the complaints or advice desks until actions 1 and 2 have reported.

Appendix A. Where things are

What	Where
This narrative and the generated column	docs/evidence/live/RUNS-AND-FINDINGS.md; docs/evidence/live/README.md
Per design: cassette, result, markdown, design, cells.json, stories	docs/evidence/live/<design>-live/
Recording timings	docs/evidence/live/timings.json
The live designs	experiments/live/
The scripts	scripts/live-designs.mjs, live-record.mjs, live-check.mjs, live-column.mjs
The plan's notes	docs/design-day2/112-REAL-ENOUGH-PLAN.md §11 (WP168, WP169)
The Spark patterns and the lease	docs/design-day2/99-DGX-SPARK.md §9
Scripted and fallible columns (full size)	docs/evidence/<design>/

Appendix B. Glossary

Term	Meaning
Desk	A bank workflow (lending, servicing, disputes, collections, onboarding, complaints, fraud, advice) run by an agent.
Book	The synthetic customers and cases a desk works through in a run; its size is the number of customers.
Design / reference design	An experiment on one desk: factors, metrics and statistics. Its id is the folder name, for example lending-stack.
Cell	One case run under one configuration of the design's factors.
Brain	What makes the bot's decisions: scripted (perfect), fallible (assumed error rate) or live (a language model).
Cartridge	A named model served by a provider, here dgx-spark/giant-qwen.
Cassette	The recorded answers of a live model for a design, from which the design replays offline with identical numbers.
Stack / policy cards	The guard control under test: policy cards placed between the bot and its actions.
Executor	Who acts on a decision: bot-everywhere, bot-with-a-person-at-the-decision, or rules-only.
Seat	A live model playing the customer across the desk (WP169).
Stories	Committed, readable renderings of sampled runs, one or more per outcome.
Wilson / Newcombe interval	95% confidence intervals for a proportion and for a difference of proportions.
Underpowered	Too few cells to detect the difference the design aimed for; "no difference" means not detectable here.
ERROR / STOPPED_BY_GUARDRAIL	A cell whose stage produced no schema-valid output; a cell halted by a guard.
Verdict	The design's conclusion about its hypothesis: supported, inconclusive, untestable or not-supported.